

Lecture 4:

Bellman Equations and Dynamic Programming

B. Ravindran

Value Functions (Recall)

The **value of a state** is the expected return when starting in state s and following π thereafter

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t | S_t = s]$$

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s\right]$$


Bellman Equation for a Policy π

$$\begin{aligned} v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) \left[r + \gamma \mathbb{E}_{\pi}[G_{t+1} | S_{t+1} = s'] \right] \end{aligned}$$

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) \left[r + \gamma v_{\pi}(s') \right]$$



- ❑ Linear equation in $|S|$ variables
- ❑ An unique solution exists

$$q_{\pi}(s, a) \doteq \sum_{s', r} p(s', r | s, a) \left[r + \gamma \sum_{a'} \pi(a' | s') q(s', a') \right]$$



An Example

- ❑ Actions: north, south, east, west (*deterministic*)
- ❑ If action would take agent off the grid: no move but reward = -1
- ❑ Other actions produce reward = 0, except actions that move agent out of special states A and B as shown.



3.3	8.8	4.4	5.3	1.5
1.5	3	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2

State-value
function
for equiprobable
random policy;
 $\gamma = 0.9$

Optimal Value Functions

- ❑ For finite MDPs, policies can be partially ordered:

$\pi \geq \pi'$ if and only if $v_\pi(s) \geq v_{\pi'}(s)$ for all $s \in \mathcal{S}$


- ❑ There is always at least one (*and possibly many*) policies that is better than or equal to all the others. This is an optimal policy. We denote them all π_*

- ❑ Optimal policies share the same optimal state-value function:


$$v_*(s) \doteq \max_{\pi} v_\pi(s) \quad \text{for all } s \in \mathcal{S}$$

$$q_*(s, a) \doteq \max_{\pi} q_\pi(s, a) \quad \text{for all } s \in \mathcal{S} \text{ and } a \in \mathcal{A}(s)$$

Example



$$M = \langle S, A, p, r \rangle$$



Many optimal policies but only one optimal value function

Bellman Optimality Equation for v_*

The value of a state under an optimal policy must equal the expected return for the best action from that state:

$$\begin{aligned} v_*(s) &= \max_{a \in \mathcal{A}(s)} q_{\pi_*}(s, a) \\ &= \max_a \mathbb{E}_{\pi_*}[G_t \mid S_t = s, A_t = a] \\ &= \max_a \mathbb{E}_{\pi_*}[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \\ &= \max_a \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')]. \end{aligned}$$

Bellman Optimality Equation for q_*


The expected return for taking action a in state s and thereafter following an optimal policy

$$\begin{aligned} q_*(s, a) &= \mathbb{E} \left[R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') \mid S_t = s, A_t = a \right] \\ &= \sum_{s', r} p(s', r \mid s, a) \left[r + \gamma \max_{a'} q_*(s', a') \right]. \end{aligned}$$

Why Optimal State-Value Functions are Useful?

- Any policy that is greedy with respect to v_* is an optimal policy.
- Therefore, given v_* , one-step-lookahead search produces the long-term optimal actions.


E.g.. back to the gridworld:



a) gridworld

0.0	-14.0	-20.0	-22.0
-14.0	-18.0	-20.0	-20.0
-20.0	-20.0	-18.0	-14.0
-22.0	-20.0	-14.0	0.0

b) V^*




c) π^*

Why Optimal Action-Value Functions are More Useful?

- Given q^* , the agent does not even have to do a one-step-ahead search.

$$\pi_*(s) = \operatorname{argmax}_{a \in A(s)} q_*(s, a)$$

$$\pi_*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$



Dynamic Programming

Dynamic Programming

- ❑ DP is the solution method of choice for MDPs
- ❑ Requires complete knowledge of system dynamics (transition matrix and rewards)
- ❑ Computationally expensive
- ❑ Curse of dimensionality
- ❑ Guaranteed to converge!

Policy Evaluation

- ❑ For a given policy π , compute the state value function v_π

- ❑ Recall Bellman equation for v_π :

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) \left[r + \gamma v_\pi(s') \right]$$

- ❑ A system of $|S|$ simultaneous linear equations
- ❑ Solve iteratively

Iterative Policy Evaluation

Input π , the policy to be evaluated

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$ arbitrarily, for $s \in \mathcal{S}$, and $V(\text{terminal})$ to 0

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$


The Bellman Operator T_π

- In the previous algo, the update to $V(s)$ can be interpreted as an operator acting on a vector V

$$T_\pi : \mathbb{R}^{|S|} \rightarrow \mathbb{R}^{|S|}$$

$$(T_\pi v)(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')]$$

Example of Policy Evaluation



Policy Improvement

- ❑ Suppose we have computed v_π for an arbitrary *deterministic policy* π
- ❑ Question: For a given state s , would it be better to choose an action $a \neq \pi(s)$?
- ❑ The value of doing a in state s is:

$$q_\pi(s, a) \doteq \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$$

- ❑ It is better to switch to action a for state s if and only if $q_\pi(s, a) > v_\pi(s)$

Policy Improvement Cont.

Do this for all states to get a new policy π' that is greedy with respect to v_π

$$\begin{aligned}\pi'(s) &= \arg \max_a q_\pi(s, a) \\ &= \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]\end{aligned}$$

Then, $v_{\pi'} \geq v_\pi$

Policy Improvement Cont.


What if $v_{\pi'} = v_{\pi}$? Then, for all $s, \in \mathcal{S}$, we have

$$v_{\pi'}(s) = \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')]$$

But this is the Bellman Optimality equation.

So $v_{\pi'} = v_*$ and both π and π' are optimal policies.

Policy Iteration



Policy Iteration Algo.

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$; $V(\text{terminal}) \doteq 0$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s', r | s, \pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)

3. Policy Improvement

policy-stable \leftarrow *true*

For each $s \in \mathcal{S}$:

old-action $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$

If *old-action* $\neq \pi(s)$, then *policy-stable* \leftarrow *false*

If *policy-stable*, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Value Iteration

- ❑ Policy evaluation step of policy iteration can be truncated without losing convergence.
- ❑ If policy evaluation step is stopped after one update of each state, we get value iteration
- ❑ Can also be interpreted as turning the Bellman optimality equation into an update rule.

$$\begin{aligned} v_{k+1}(s) &\doteq \max_a \mathbb{E}[R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) \left[r + \gamma v_k(s') \right], \end{aligned}$$

Value iteration Algo.

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

```
|  $\Delta \leftarrow 0$   
| Loop for each  $s \in \mathcal{S}$ :  
|    $v \leftarrow V(s)$   
|    $V(s) \leftarrow \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$   
|    $\Delta \leftarrow \max(\Delta, |v - V(s)|)$   
until  $\Delta < \theta$ 
```

Output a deterministic policy, $\pi \approx \pi_*$, such that


$$\pi(s) = \arg\max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$$

Asynchronous DP

- ❑ Disadvantage of algorithms discussed is we have to do the updates over the entire state set
- ❑ In asynchronous DP, the updates are not done over the entire state set at each iteration
- ❑ Have to ensure that every state is visited sufficiently often for convergence
- ❑ Gives flexibility to choose order of updates
- ❑ Can intertwine real time interaction with the environment and DP updates
- ❑ Can focus updates on parts of state space relevant to agent

Real-Time DP (RTDP)

- ❑ On-policy trajectory-sampling version of value-iteration algorithm.
- ❑ Updates values of states visited in the actual trajectory




1. Take action according to π
2. Update $V_{\pi}(s)$
3. Update $\pi(a|s)$

- ❑ Unlike asynchronous-DP, no requirement to update every state infinitely often.

Generalized Policy Iteration

- ❑ GPI refers to the idea of letting policy evaluation and policy improvement interact, independent of their granularity.



GPI

- ❑ Almost all RL methods can be viewed as GPI.
- ❑ Policy iteration has evaluation running to completion before improvement begins.
- ❑ In value iteration, only one step of evaluation is done before the improvement step.
- ❑ In Asynchronous DP, the two are interleaved at a finer granularity.

