


Lecture 5: Temporal Difference Learning and Monte-Carlo Methods

B. Ravindran


Dynamic Programming

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) \left[r + \gamma v_{\pi}(s') \right]$$




Monte-Carlo RL

$$v_\pi(s) \doteq \mathbb{E}_\pi[G_t \mid S_t = s]$$




Best of Both Worlds !




Simplest 'TD' Method


$$v_{\pi}(s) = R_{t+1} + \gamma v_{\pi}(s')$$




DP vs MC vs TD



DP



MC



TD

Bootstrapping and Sampling


❏ Bootstrapping: Update using an estimate

- ❏ DP and TD bootstrap ✓
- ❏ Monte Carlo does not bootstrap ✗

❏ Sampling: Update using samples

- ❏ TD and Monte Carlo sample ✓
- ❏ DP (*typically*) does not sample ✗

Bootstrapping and Sampling



Monte-Carlo RL

- ❑ Learning directly from sample episodes of experience
- ❑ Does not use a known model and is model-free
- ❑ MC does not use bootstrapping
- ❑ Value functions are calculated as mean of discounted returns (G_t)

Monte-Carlo Prediction

- ❑ First-visit MC method estimates $V(s)$ as the average of the returns following first visits to s
- ❑ Every-visit MC method estimates $V(s)$ as the average of the returns following all visits to s

Monte-Carlo Prediction: First Visit

Input: a policy π to be evaluated

Initialize:

$V(s) \in \mathbb{R}$, arbitrarily, for all $s \in \mathcal{S}$

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Loop forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$


Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless S_t appears in S_0, S_1, \dots, S_{t-1} :

Append G to $Returns(S_t)$

$V(S_t) \leftarrow \text{average}(Returns(S_t))$



Monte-Carlo Control: First Visit

Algorithm parameter: small $\varepsilon > 0$

Initialize:

$\pi \leftarrow$ an arbitrary ε -soft policy

$Q(s, a) \in \mathbb{R}$ (arbitrarily), for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$ empty list, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair S_t, A_t appears in $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$:


Append G to $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \operatorname{argmax}_a Q(S_t, a)$ (with ties broken arbitrarily)

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$



Temporal Difference Learning

Temporal Difference

- ❑ *“If one had to identify one idea as central and novel to RL, it would undoubtedly be TD learning”*

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$



- ❑ Simple rule to explain complex behaviors
- ❑ Has had profound impact in behavioral psychology and neuroscience!

Temporal Difference

- ❑ TD methods do not require a model of the environment, only experience
- ❑ TD methods can be fully incremental (bootstrapping)
 - ❑ You can learn **before** knowing the final outcome
 - ❑ Less memory & peak computation
 - ❑ You can learn **without** the final outcome
 - ❑ From incomplete sequences
- ❑ TD methods thus combine individual advantages of DP and MC


TD Prediction

- ❑ Policy Evaluation (the prediction problem): for a given policy, compute the state-value function
- ❑ No knowledge of p & r , but access to the real system, or a “sample” model assumed

- ❑ Uses bootstrapping and sampling

- ❑ The simplest TD method, TD(0):

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$



Stochastic Averaging Rule

$$E(x) \approx \frac{1}{n} \sum_{i=1}^n x_i$$


Let \bar{x}_n be the average of the first n samples

$$\begin{aligned}\bar{x}_{n+1} &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n) \\ &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n + \bar{x}_n - \bar{x}_n) \\ &= \frac{1}{n+1} ((n+1)\bar{x}_n + (x_{n+1} - \bar{x}_n)) \\ &= \bar{x}_n + \frac{1}{n+1} (x_{n+1} - \bar{x}_n) \\ &= \bar{x}_n + \alpha (x_{n+1} - \bar{x}_n)\end{aligned}$$

new estimate = old estimate + α (new sample - old estimate)

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

TD Update Example



Assuming :

reward, r , $A \rightarrow B : 0$

$\alpha : 0.2$


$\gamma : 0.9$

$$V(A) = V(A) + \alpha[r + \gamma V(B) - V(A)]$$

$$V(A) = 0.5 + 0.2[0 + 0.9 * 0.8 - 0.5] = 0.544$$

TD Control

- ❑ The control problem: approximate optimal policies
- ❑ Recall the idea of GPI:



- ❑ Evaluation: use TD(0) to evaluate value function
- ❑ Improvement: make policy **greedy** wrt current value function
- ❑ Note that we estimate action values rather than state values in the absence of a model

ϵ -Greedy Policies

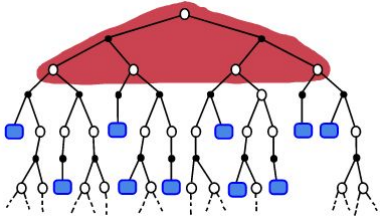


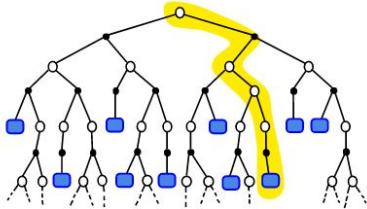


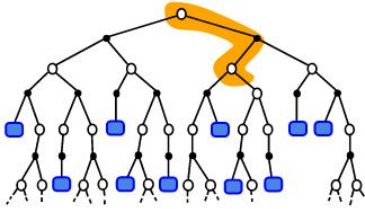


$$A^* \leftarrow \operatorname{argmax}_a Q(S_t, a)$$

For all $a \in \mathcal{A}(S_t)$:

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \epsilon + \epsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \epsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

- Any ϵ -greedy policy with respect to Q following π is an improvement over any ϵ -soft policy π is assured by the policy improvement theorem


Summary: DP vs MC vs TD

		Bootstrapping?	Sampling?	Bias/Variance
DP				-
MC				Low Bias, High Variance
TD				High Bias, Low Variance

Sarsa

Bellman Equation

$$q_{\pi}(s, a) \doteq \sum_{s', r} p(s', r | s, a) [r + \gamma \sum_{a'} \pi(a' | s') q(s', a')]$$




$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Temporal Difference
Error

Sarsa: On-Policy TD Control

- ❑ On-policy control: Improve the behaviour policy

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$



- ❑ Convergence is guaranteed as long as it is GLIE (Greedy in the Limit with Infinite Exploration)
 - ❑ all state-action pairs are visited an ∞ no. of times
 - ❑ the policy converges in the limit to the greedy policy

Sarsa: On-Policy TD Control

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ε -greedy)

Loop for each step of episode:


Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ε -greedy) IMPROVEMENT


$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$ EVALUATION

$S \leftarrow S'; A \leftarrow A';$

until S is terminal



Sarsa: On-Policy TD Control



Every **time-step**:


Policy evaluation **Sarsa**, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

Q-Learning

Bellman Optimality Equation

$$q_*(s, a) = \mathbb{E} \left[R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') \mid S_t = s, A_t = a \right]$$



$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Temporal Difference
Error

Q-Learning: Off-Policy TD Control

- ❑ In off-policy control, we have two policies:
 - ❑ behavior policy – used to generate behavior
 - ❑ estimation policy – the policy that is being evaluated and improved

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Q-Learning: Off-Policy TD Control

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:


 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R , S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal



Q-Learning: Off-Policy TD Control

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Loop for each step of episode:


Choose A from S using policy derived from Q (ε -greedy)

Take action A , observe R , S'


$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$


until S is terminal



Cliff Walking: SARSA vs Q-learning



- ❑ $R = -1$ on all transitions except those into the region marked "The Cliff"
- ❑ Stepping into this region incurs a reward of -100 and the agent restarts





n - step Bootstrapping

n -Step TD Prediction

- ❑ Consider TD(0): $V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$
- ❑ Here, the target (for estimating TD error) contains only the next step reward: $G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$
- ❑ Alternatively, we can consider the rewards received in the next n steps:

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$$

- ❑ The extreme would be to consider rewards till the end of the episode (Monte Carlo)

n -Step TD Prediction

1-step TD
and TD(0)



2-step TD



3-step TD




...

n -step TD




...

∞ -step TD
and Monte Carlo




The spectrum of
back-ups from
one-step TD to
up-until-termination MC


n -Step TD Prediction - Example




Average
RMS error
over 19 states
and first 10
episodes



n -Step TD Prediction - Example



Average
RMS error
over 19 states
and first 10
episodes



But what value of n to choose?

Why not average all n -step backups?

From n -Step TD to TD(λ)

- ❑ Instead of using *1 n -step backup*, we can consider an average of *multiple n -step backups*

- ❑ Example: $G_t^{avg} = \frac{1}{2} G_{t:t+10} + \frac{1}{2} G_{t:t+20}$

estimates of the same value - $V(s_t)$


- ❑ In TD(λ), the average contains all the n -step backups each weighted proportional to λ^{n-1} , where $0 \leq \lambda \leq 1$.

TD(λ)


- In TD(λ), the average contains all the n-step backups each weighted proportional to λ^{n-1} , where $0 \leq \lambda \leq 1$.

- λ -return:

$$G_t^\lambda \doteq (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n}$$




TD(λ) vs n-step TD



- ❑ The results with the TD(λ) are slightly better at the best values of α and λ , and at high α


TD(λ) - Forward View

- ❑ TD(λ) as a *forward* - view algorithm
- ❑ Wait till the episode terminates
- ❑ Then for each state visited, we look *forward* in time to all the future rewards and combine them



Eligibility Traces

- ❑ To implement TD(λ) on-the-fly we use the concept of eligibility traces
- ❑ These are variables associated with each state denoted by $E_t(s)$
- ❑ They indicate the degree to which each state is eligible for undergoing learning changes
- ❑ On each step:
$$E_t(s) = \begin{cases} \gamma\lambda E_{t-1}(s) & \text{if } s \neq S_t; \\ \gamma\lambda E_{t-1}(s) + 1 & \text{if } s = S_t, \end{cases}$$



TD(λ) using Eligibility Traces

Initialize $V(s)$ arbitrarily (but set to 0 if s is terminal)

Repeat (for each episode):


Initialize $E(s) = 0$, for all $s \in \mathcal{S}$

Initialize S

Repeat (for each step of episode):


$A \leftarrow$ action given by π for S

Take action A , observe reward, R , and next state, S'

 $\delta \leftarrow R + \gamma V(S') - V(S)$

$E(S) \leftarrow E(S) + 1$ (accumulating traces)

For all $s \in \mathcal{S}$:

 $V(s) \leftarrow V(s) + \alpha \delta E(s)$

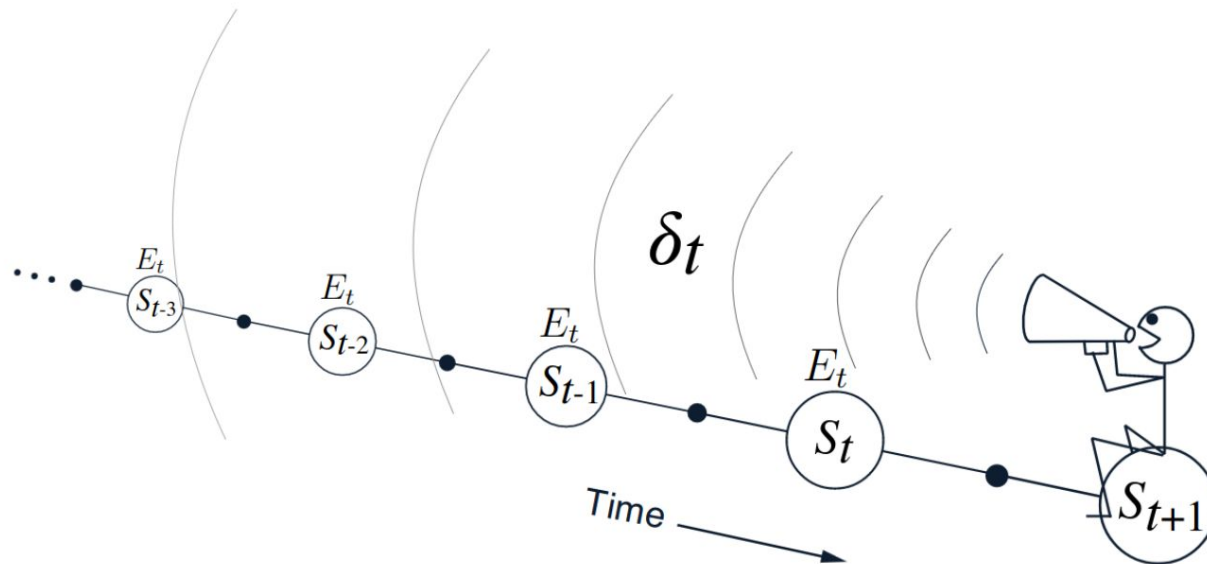
$E(s) \leftarrow \gamma \lambda E(s)$

$S \leftarrow S'$

until S is terminal


TD(λ) - Backward View

- ❑ TD(λ) as a *backward* - view algorithm
- ❑ At each moment the current TD error is assigned it *backwards* to each prior state according to how eligible it is
- ❑ Need not wait for the episode to terminate!




Speedup in Policy Learning


Path taken




Action values increased by one-step Sarsa



Action values increased by 10-step Sarsa



Action values increased by Sarsa(λ) with $\lambda=0.9$



TD(0) vs TD(λ)



Example grid world domain


Grey cells = obstacles

Notice $R = 1.0$ for centre cell with most surrounding cells having $R = -1.0$

Cell values = value function estimates

TD(0) vs TD(λ)

0.12 →	0.14 →	0.17 →	0.20 ↓	0.22 ↓	0.22 ↓	-0.00 ↕	-0.04 ↔	-0.03 ↕	-0.03 ↕
0.10 →	0.16 →	0.20 →	0.23 →	0.27 →	0.31 ↓	0.21 ↔	-0.04 ↔	-0.03 ↕	-0.03 ↕
-0.05 ↓					0.36 ↓				-0.02 ↕
-0.05 ↕	-0.05 ↕	-0.05 ↕	-0.77 ↕ R -1.0		0.41 →	0.47 →	0.53 ↓	0.37 ←	0.03 ↕
-0.05 ↕	-0.05 ↕	-0.04 ↔	-0.04 ↕		-0.07 ↓ R -1.0	-0.48 → R -1.0	0.60 ↓	0.51 ←	0.25 ↔
-0.04 ↔	-0.04 ↕	-0.04 ↓	-0.04 ↕		1.09 ← R 1.0	-0.04 ← R -1.0	0.68 ↓	-0.44 ↔ R -1.0	0.05 ↕
-0.04 ↕	-0.04 ↕	-0.04 ↕	-0.04 ↔		0.97 ↕	0.87 ← R -1.0	0.77 ←	-0.39 ↔ R -1.0	-0.01 ↕
-0.04 ↕	-0.04 ↕	-0.03 ↓	-0.68 ↕ R -1.0		-0.20 ↕ R -1.0	-0.26 ↕ R -1.0	0.67 ↕	0.34 ↔	0.05 ↔
-0.04 ↓	-0.03 ↔	-0.00 ↔	0.05 →	0.16 →	0.29 →	0.42 →	0.53 ↕	0.21 ↔	-0.01 ↕
-0.04 ↔	-0.04 ↕	-0.03 ↕	-0.02 ↕	-0.01 ↕	0.03 ↕	0.10 ↕	0.11 ↕	-0.01 ↔	-0.01 ↕




Algo : Q - learning

$\gamma = 0.9, \varepsilon = 0.2, \alpha = 0.1, \lambda = 0$



TD(0) vs TD(λ)

0.20 ↓	0.23 ↓	0.22 ↓	0.25 ↓	0.29 ↓	0.33 ↓	0.12 ↓	-0.02 ↕	-0.03 ↻	-0.03 ↻
0.23 →	0.25 →	0.28 →	0.31 →	0.34 →	0.38 ↓	0.33 ←	0.02 ↔	-0.03 ↕	-0.03 ↕
0.20 ↕					0.41 ↓				-0.01 ↕
0.03 ↕	-0.04 ↔	-0.04 ↔	-0.48 ↕ R -1.0		0.45 →	0.49 →	0.54 ↓	0.47 ←	0.06 ↕
-0.04 ↕	-0.04 ↕	-0.04 ↕	-0.04 ↔		0.13 ↓ R -1.0	-0.17 → R -1.0	0.59 ↓	0.52 ←	0.22 ←
-0.04 ↕	-0.04 ↔	-0.04 ↕	-0.04 ↕		0.95 → R 1.0	0.14 ← R -1.0	0.65 ↓	-0.19 ← R -1.0	0.05 ↕
-0.04 ↕	-0.04 ↕	-0.04 ↕	-0.04 ↔		0.87 ↕	0.78 ←	0.71 ←	-0.12 ← R -1.0	-0.00 ↕
-0.04 ↕	-0.04 ↕	-0.04 ↔	-0.48 ↕ R -1.0		0.09 ↕ R -1.0	-0.02 ↕ R -1.0	0.62 ↕	0.38 ←	0.12 ↔
-0.04 ↔	-0.04 ↕	-0.04 ↔	-0.03 ↔	-0.02 ↔	0.02 ↔	0.09 →	0.26 ↕	0.02 ↕	-0.02 ↔
-0.04 ↕	-0.04 ↔	-0.04 ↕	-0.04 ↕	-0.03 ↔	-0.03 ↔	-0.02 ↕	-0.02 ↔	-0.02 ↕	-0.02 ↔



Algo : Q - learning

$\gamma = 0.9, \varepsilon = 0.2, \alpha = 0.1, \lambda = 0.2$





Importance Sampling

More on Off-Policy Learning

- ❑ Evaluate target policy $\pi(a/s)$ to compute $v_\pi(s)$ or $q_\pi(s, a)$
- ❑ Follow behaviour policy $\mu(a/s)$
- ❑ Assumption of coverage:
 - ❑ Every action taken under π is also taken, at least occasionally, under μ
 - ❑ $\pi(a/s) > 0$ implies $\mu(a/s) > 0$
- ❑ e.g., π can be greedy while μ can be ε -greedy

Importance Sampling

- A general technique for estimating expected values under one distribution given samples from another.

$$\begin{aligned}\mathbb{E}_{X \sim P}[f(X)] &= \sum P(X)f(X) \\ &= \sum Q(X) \frac{P(X)}{Q(X)} f(X) \\ &= \mathbb{E}_{X \sim Q} \left[\frac{P(X)}{Q(X)} f(X) \right]\end{aligned}$$

Off-Policy MC with Weighted Importance Sampling

Importance-sampling ratio:

$$\rho_t^T = \frac{\prod_{k=t}^{T-1} \pi(A_k|S_k)p(S_{k+1}|S_k, A_k)}{\prod_{k=t}^{T-1} \mu(A_k|S_k)p(S_{k+1}|S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{\mu(A_k|S_k)}.$$

Weighted average return for V :

$$V(s) \doteq \frac{\sum_{t \in \mathcal{T}(s)} \rho_{t:T(t)-1} G_t}{|\mathcal{T}(s)|}$$