

Lecture 7:

Function Approximation, SGD, DQN


B. Ravindran

Need for Function Approximation

- ❑ Issues with large state/action spaces:
 - ❑ tabular approaches not memory efficient
 - ❑ data sparsity
 - ❑ continuous state/action spaces
 - ❑ generalization
- ❑ Use a parameterized representation
 - ❑ Value Functions
 - ❑ Policies
 - ❑ Models

Non-Linear Function Approximator

- ❑ Linear function approximators are restrictive. Can only model linear functions
- ❑ Basis expansion does help to generate non-linear functions in the original input space
- ❑ Non-linear approximators can model complex functions and are very powerful
- ❑ The features are learnt on the fly and are not hard coded as is the case with tile and sparse coding
- ❑ Can generalize to unseen states
- ❑ **Disadvantage: Requires a lot of data and compute**



Gradient Descent

- ❑ Gradient Descent- first-order iterative optimization algorithm for finding a local minimum of a differentiable function
- ❑ Compute the gradient of the function at the current point and take a step in the opposite direction. This is the direction of steepest descent
- ❑ The logic behind Gradient Descent can be understood by considering the Taylor Series expansion of the function around the current point, up to first order

$$f(x_0 + h) \approx f(x_0) + h \cdot f'(x_0)$$

$$L(\theta + \alpha \Delta \theta) \approx L(\theta) + \alpha \Delta \theta^T \nabla_{\theta} L(\theta)$$

Semi Gradient Methods

- ❑ While computing the gradient of the TD error in Q-learning, we typically ignore the gradient of the TD target
- ❑ Hence, it is a Semi Gradient method i.e we are computing an approximation of the true gradient

$$Q(s_t, a_t) = \phi^T(s_t, a_t) \times w_t$$


$$\delta_t = r_{t+1} + \gamma \max_a Q(s_t, a) - Q(s_t, a_t) \quad \text{TD-ERROR}$$

$$\nabla_{w_t} \left[r_{t+1} + \gamma \max_a Q(s_t, a) - Q(s_t, a_t) \right]^2 = -\delta_t \phi(s_t, a_t)$$

$$w_{t+1} = w_t + \alpha \delta_t \phi(s_t, a_t)$$


Stochastic Gradient Descent

- ❑ In Stochastic Gradient Descent, the true gradient of the loss function is approximated by the gradient at a single example.
- ❑ In practice, we usually perform Mini Batch Gradient Descent, where we compute the gradient using a mini batch of examples. This allows for more efficient computation and smoother convergence.




TD-Gammon

- ❑ TD-Gammon (Tesauro 92, 94, 95)
- ❑ Beat the best human player in 1995
- ❑ Learnt completely by self play
- ❑ New moves not recorded by humans in centuries of play




TD-Gammon

TD-Gammon (Tesauro 92, 94,



New moves not recorded by humans in centuries of play





DQN

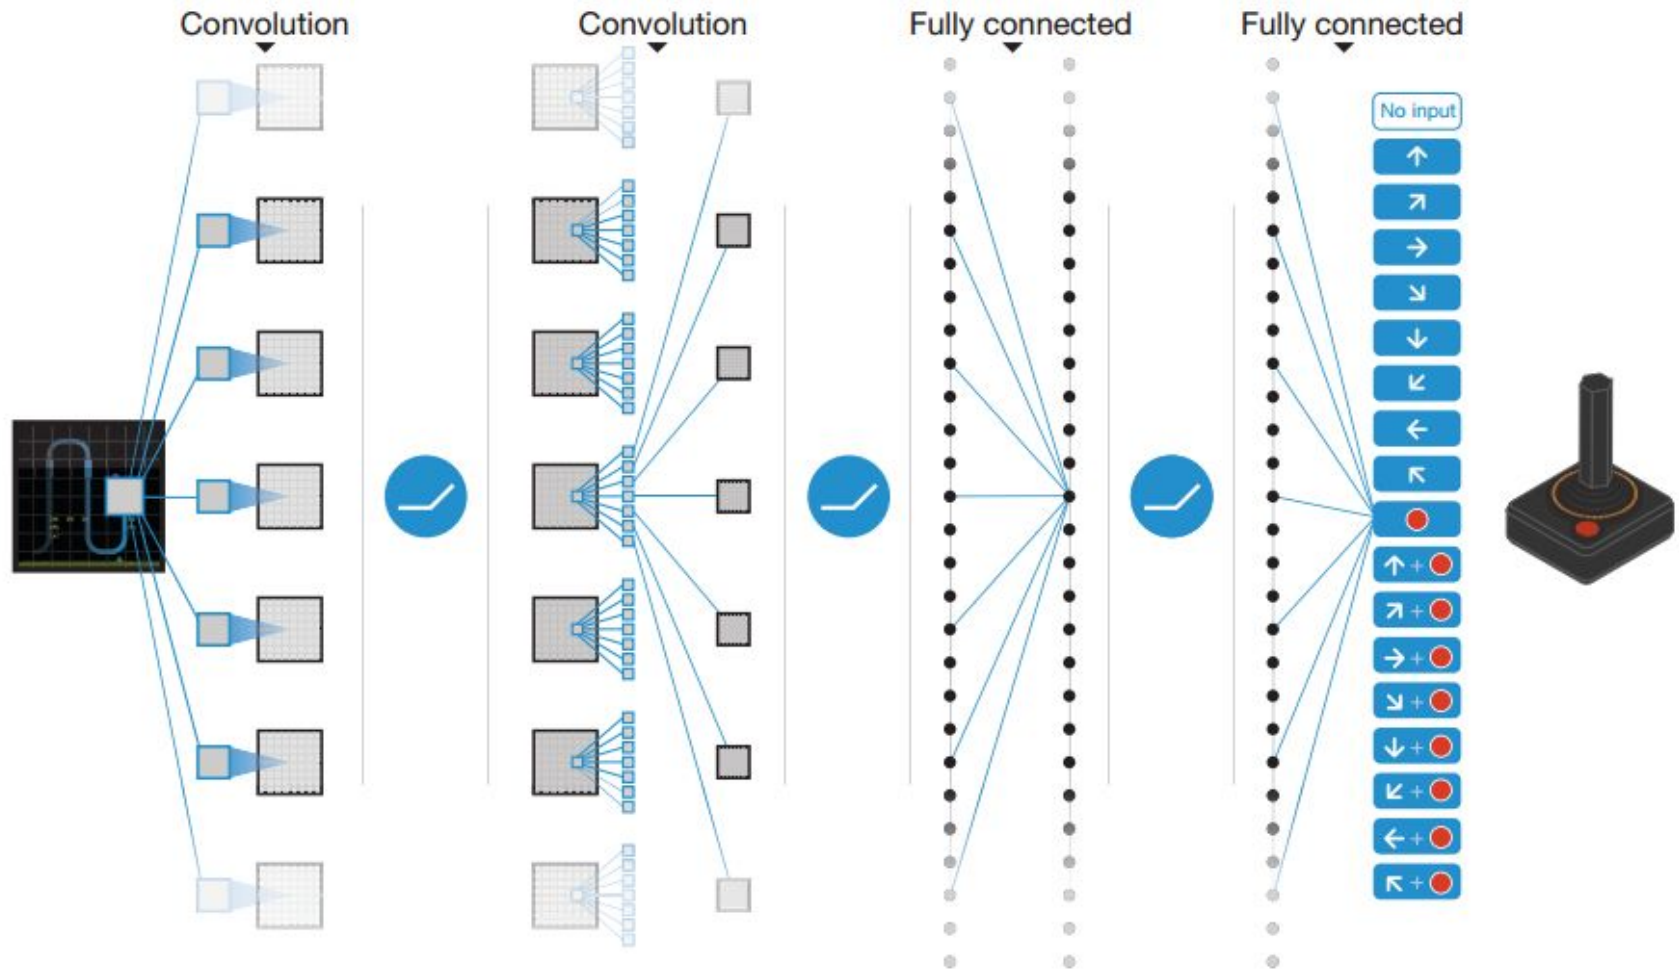
Deep Q-Learning Network (DQN)



Q. What about input features?

A. Learnt to play from video input from scratch!

DQN



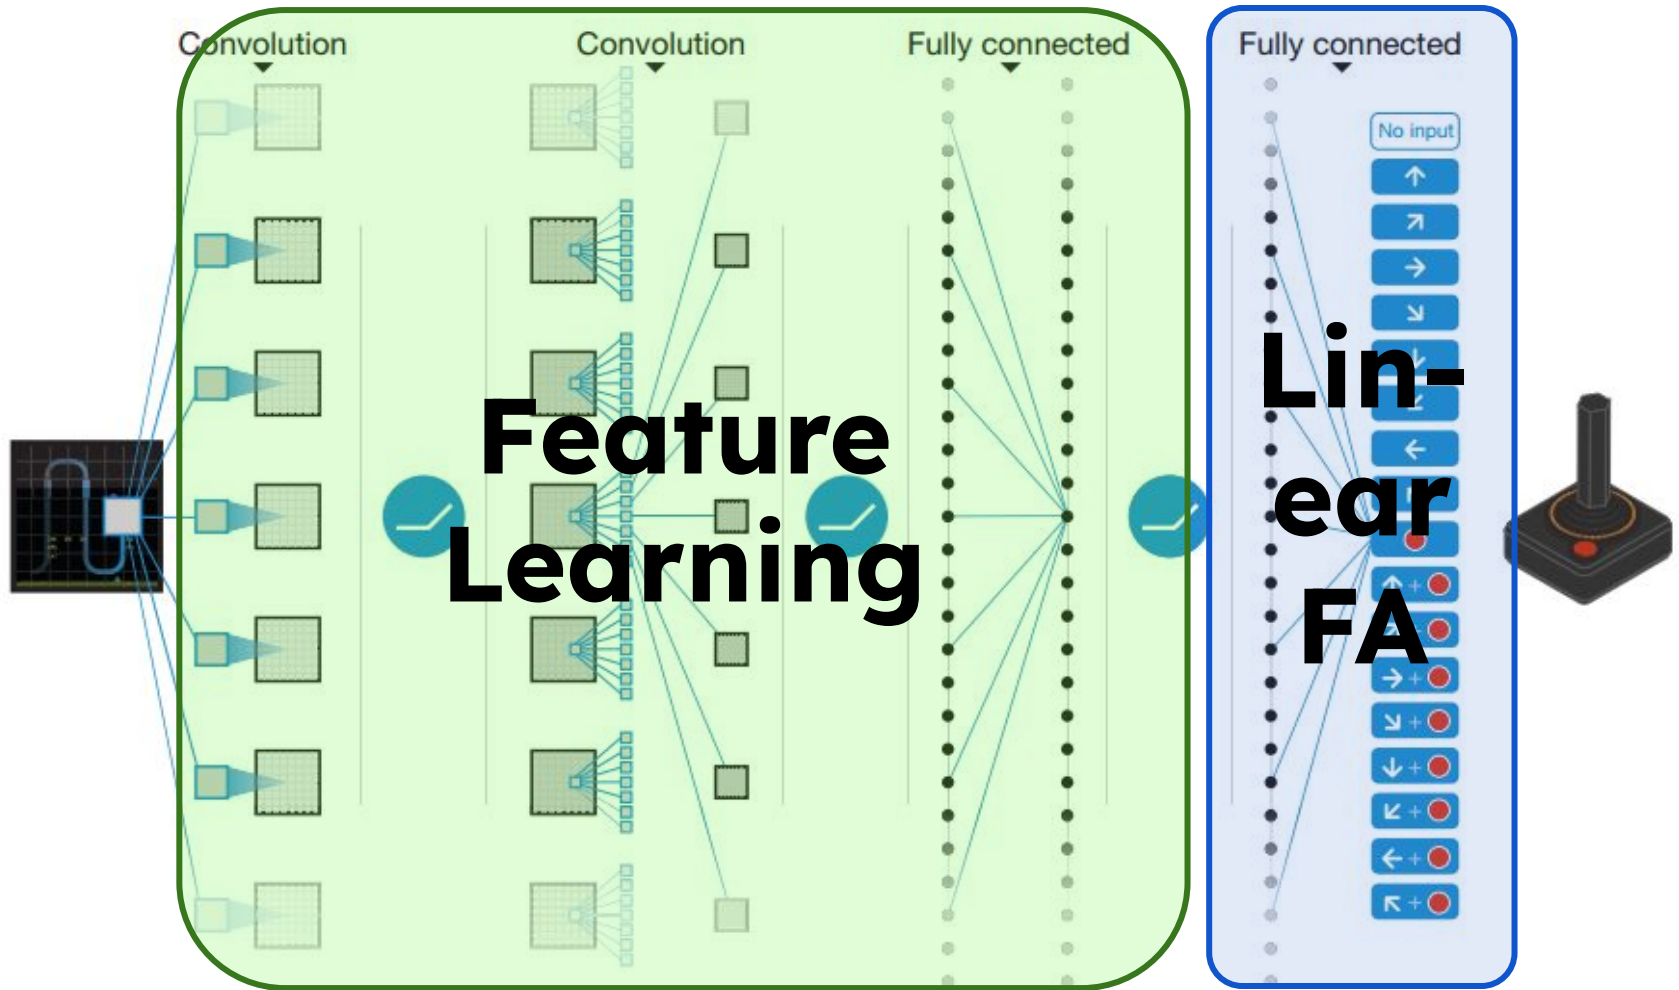
Source: Deep Q Networks, Nature 2015

DQN

- ❑ Input: $84 \times 84 \times 4$
- ❑ Layer 1: conv 8×8 , stride=4, 16 filters -- ReLU
- ❑ Layer 2: conv 4×4 , stride=2, 32 filters -- ReLU
- ❑ Layer 3: fully_connected 256 - ReLU
- ❑ Output layer: fully_connected 18 - Linear

Source: Deep Q Networks, Nature 2015

DQN




Source: Deep Q Networks, Nature 2015

Q-Network Learning


$$w_{t+1} = w_t - \frac{1}{2}\alpha \nabla_{w_t} \left[r_{t+1} + \gamma \max_a \hat{q}(s_{t+1}, a) - \hat{q}(s_t, a_t) \right]^2$$

- ❑ **Divergence** is an issue since the current network is used to decide its own target
- ❑ Correlations between samples
- ❑ Non-stationarity of the targets
- ❑ How do we address these issues?
 - ❑ Replay Memory
 - ❑ Freeze target network



Q-Network Learning

- ❑ **Replay Memory/Buffer** : To remove correlations, we build data-set from the agent's experience




- ❑ **Frozen Target Network** : Sample experiences from dataset; w^- frozen (with periodic updates) to address non-stationarity

$$\left[\left\{ r_{t+1} + \gamma \max_a \hat{q}(s_{t+1}, a; w^-) \right\} - \hat{q}(s_t, a_t; w) \right]^2$$


DQN

- ❑ Architecture:
 - ❑ Has a set of convolutional layers which act as feature extractors
 - ❑ These features are then passed through a series of fully connected layers
 - ❑ The output layer has $|A|$ number of nodes which are used to calculate the Q-value for each action
- ❑ The network is updated using huber loss and not regular least squares loss



$$L_{\delta}(a) = \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta, \\ \delta(|a| - \frac{1}{2}\delta), & \text{otherwise.} \end{cases}$$





DQN - Design Choices



Type 1 : Fully Sharing



Type 2 : Semi-Sharing



Type 3 : No sharing