

Lecture 9: Actor Critic Methods

B. Ravindran

Recall - REINFORCE: MC PG

$$\begin{aligned}\nabla J(\boldsymbol{\theta}) &\propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \boldsymbol{\theta}) \\&= \mathbb{E}_\pi \left[\sum_a q_\pi(S_t, a) \nabla \pi(a|S_t, \boldsymbol{\theta}) \right] && \text{Expectation over the visited states while following } \Pi \\&= \mathbb{E}_\pi \left[q_\pi(S_t, A_t) \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right] && \begin{array}{l} \text{Multiply and divide by } \Pi(a|s, \boldsymbol{\theta}) + \\ \text{Expectation over the actions taken while following } \Pi \end{array} \\&= \mathbb{E}_\pi \left[G_t \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right]\end{aligned}$$

The gradient update is

- proportional to the return
- inversely proportional to the action probability

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha G_t \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta}_t)}{\pi(A_t|S_t, \boldsymbol{\theta}_t)}$$

REINFORCE w/ Baseline

- ❑ The policy gradient theorem can be generalized to include a comparison of the action value to an arbitrary baseline $b(s)$

$$\nabla J(\boldsymbol{\theta}) \propto \sum_s \mu(s) \sum_a \left(q_\pi(s, a) - b(s) \right) \nabla \pi(a|s, \boldsymbol{\theta})$$

- ❑ Baseline - should be a function that is independent of the action a

$$\sum_a b(s) \nabla \pi(a|s, \boldsymbol{\theta}) = b(s) \nabla \sum_a \pi(a|s, \boldsymbol{\theta}) = b(s) \nabla 1 = 0$$

- ❑ Update rule of REINFORCE with baseline

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha \left(G_t - b(S_t) \right) \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta}_t)}{\pi(A_t|S_t, \boldsymbol{\theta}_t)}$$

(Doesn't change the expected value, but has an effect on the variance)

Recall - Actor-Critic Methods

- ❑ Actor-Critic methods learn both a policy and a state-value function simultaneously
- ❑ The policy is referred to as the actor that suggests actions given a state
- ❑ The value function is referred to as the critic. It evaluates actions taken by the actor based on the given policy

$$\begin{aligned}\boldsymbol{\theta}_{t+1} &\doteq \boldsymbol{\theta}_t + \alpha (G_{t:t+1} - \hat{v}(S_t, \mathbf{w})) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)} \\ &= \boldsymbol{\theta}_t + \alpha (R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)} \\ &= \boldsymbol{\theta}_t + \alpha \delta_t \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)}\end{aligned}$$

Recall - One-step Actor-Critic

Input: a differentiable policy parameterization $\pi(a|s, \boldsymbol{\theta})$
Input: a differentiable state-value function parameterization $\hat{v}(s, \mathbf{w})$
Parameters: step sizes $\alpha^{\boldsymbol{\theta}} > 0$, $\alpha^{\mathbf{w}} > 0$
Initialize policy parameter $\boldsymbol{\theta} \in \mathbb{R}^{d'}$ and state-value weights $\mathbf{w} \in \mathbb{R}^d$ (e.g., to $\mathbf{0}$)
Loop forever (for each episode):
 Initialize S (first state of episode)
 $I \leftarrow 1$
 Loop while S is not terminal (for each time step):
 $A \sim \pi(\cdot|S, \boldsymbol{\theta})$
 Take action A , observe S', R
 $\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$ (if S' is terminal, then $\hat{v}(S', \mathbf{w}) \doteq 0$)
 $\mathbf{w} \leftarrow \mathbf{w} + \alpha^{\mathbf{w}} \delta \nabla \hat{v}(S, \mathbf{w})$
 $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha^{\boldsymbol{\theta}} I \delta \nabla \ln \pi(A|S, \boldsymbol{\theta})$
 $I \leftarrow \gamma I$
 $S \leftarrow S'$

(This is a fully online, incremental algorithm, with states, actions, and rewards processed as they occur and then never revisited again)

Comparison to REINFORCE

REINFORCE w/ baseline

$$\theta_{t+1} \doteq \theta_t + \alpha \left(G_t - b(S_t) \right) \frac{\nabla \pi(A_t | S_t, \theta_t)}{\pi(A_t | S_t, \theta_t)}$$

- ❑ G_t - unbiased estimate but high variance
- ❑ Recall $\mathbb{E}_\pi[G_t | S_t, A_t] = q_\pi(S_t, A_t)$
- ❑ Need a good estimate of $q_\pi(S_t, A_t)$ with less variance than G_t
- ❑ In 1-step Actor-Critic, we use \hat{v} for both estimating $q_\pi(S_t, A_t)$ and as the baseline
- ❑ This introduces bias, but lesser variance - Leads to better performance


1-step Actor-Critic

$$\theta_{t+1} \doteq \theta_t + \alpha \left(R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w}) \right) \frac{\nabla \pi(A_t | S_t, \theta_t)}{\pi(A_t | S_t, \theta_t)}$$

Common Features of AC Methods


❑ Actor: Computes the policy π_θ and updates θ

❑ Critic: Computes an estimate $\hat{v}(s, w)$ of the state value function. Updates the parameter w



Basic Actor Critic Algorithm

1. Take action $\mathbf{a} \sim \pi_{\theta}(\mathbf{a} \mid \mathbf{s})$ and receive $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$
2. Update value parameter \mathbf{w} using data $(\mathbf{s}, r + \gamma \hat{v}(\mathbf{s}', \mathbf{w}))$
3. Compute $\hat{\delta}(\mathbf{s}, \mathbf{a}) = r + \gamma \hat{v}(\mathbf{s}', \mathbf{w}) - \hat{v}(\mathbf{s}, \mathbf{w})$
4. $\theta \leftarrow \theta + \alpha \cdot \hat{\delta}(\mathbf{s}, \mathbf{a}) \cdot \nabla_{\theta} \log \pi_{\theta}(\mathbf{a} \mid \mathbf{s})$



Critic Update

1. Take action $\mathbf{a} \sim \pi_{\theta}(\mathbf{a} \mid \mathbf{s})$ and receive $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$
2. Update value parameter \mathbf{w} using data $(\mathbf{s}, r + \gamma \hat{v}(\mathbf{s}', \mathbf{w}))$
3. Compute $\hat{\delta}(\mathbf{s}, \mathbf{a}) = r + \gamma \hat{v}(\mathbf{s}', \mathbf{w}) - \hat{v}(\mathbf{s}, \mathbf{w})$
4. $\theta \leftarrow \theta + \alpha \cdot \hat{\delta}(\mathbf{s}, \mathbf{a}) \cdot \nabla_{\theta} \log \pi_{\theta}(\mathbf{a} \mid \mathbf{s})$


- ❑ Step 2 of the previous algorithm usually happens in batches
- ❑ Minimize the squared loss:

$$L(\mathbf{w}) = \frac{1}{N} \sum_i \|\hat{v}(\mathbf{s}_i, \mathbf{w}) - y_i\|^2$$


Batch Size

Design Choices

Two Network Design



Shared Network Design






Advantage Actor-Critic Methods

Recall - Advantage Function

- ❑ The advantage function is the difference between the q-value and the value function
- ❑ It can be interpreted as a measure of the advantage of taking action a in state s as compared to following policy π

$$A_{\pi}(\mathbf{s}, \mathbf{a}) = q_{\pi}(\mathbf{s}, \mathbf{a}) - v_{\pi}(\mathbf{s})$$

A3C – Asynchronous Advantage Actor Critic



A3C Algorithm

// Assume global shared parameter vectors θ and θ_v and global shared counter $T = 0$

// Assume thread-specific parameter vectors θ' and θ'_v

Initialize thread step counter $t \leftarrow 1$

repeat

Reset gradients: $d\theta \leftarrow 0$ and $d\theta_v \leftarrow 0$.

Synchronize thread-specific parameters $\theta' = \theta$ and $\theta'_v = \theta_v$

$t_{start} = t$

Get state s_t

repeat

Perform a_t according to policy $\pi(a_t|s_t; \theta')$

Receive reward r_t and new state s_{t+1}

$t \leftarrow t + 1$

$T \leftarrow T + 1$

until terminal s_t **or** $t - t_{start} == t_{max}$

$$R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{Bootstrap from last state} \end{cases}$$

for $i \in \{t - 1, \dots, t_{start}\}$ **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt θ' : $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$

Accumulate gradients wrt θ'_v : $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

end for

Perform asynchronous update of θ using $d\theta$ and of θ_v using $d\theta_v$.

until $T > T_{max}$


Reset thread
params,
update local
params with
global params

Gather
experience

Compute the
gradients for
this thread

Update
global
params


A2C – Synchronous Advantage Actor Critic




A3C vs A2C

- ❑ We remove the "asynchronous" part of A3C
- ❑ The updates to the global parameters are executed only after all the threads have finished their computation

asynchronous parallel actor-critic



synchronized parallel actor-critic



Compatible Parametrization

- ❑ Substituting the approximation $\hat{q}(s, a, w)$ instead of the true value of $q_\pi(s, a)$ may introduce bias
- ❑ It can be proved that there is no bias if the function approximator has a “compatible” parametrization with the policy parametrization

- ❑ Condition 1: projection of characteristic eligibility

$$\hat{q}(\mathbf{s}, \mathbf{a}, \mathbf{w}) = \nabla_\theta \log \pi_\theta(\mathbf{a} \mid \mathbf{s})^T \mathbf{w}$$

- ❑ Condition 2: minimize mean squared error

$$\mathbf{w} = \arg \min \mathbb{E}_{\mathbf{s} \sim \rho^{\pi_\theta}, \mathbf{a} \sim \pi_\theta} \left[(\hat{q}(\mathbf{s}, \mathbf{a}, \mathbf{w}) - q_{\pi_\theta}(\mathbf{s}, \mathbf{a}))^2 \right]$$