

Lecture 10:

(Deep) Deterministic Policy Gradient

B. Ravindran

PG for Continuous Action Spaces

❑ Policy Gradient Theorem (Discrete Actions):

$$\begin{aligned}\nabla J(\theta) &\propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a \mid s) \\ &= \mathbb{E}_{s \sim \mu^\pi(s), a \sim \pi_\theta} [q_\pi(S_t, A_t) \nabla \log \pi(A_t \mid S_t)]\end{aligned}$$

❑ Policy Gradient Theorem (Continuous Actions):

$$\begin{aligned}\nabla J(\theta) &= \int_{\mathcal{S}} \mu^\pi(s) \int_{\mathcal{A}} q_\pi(s, a) \nabla \pi(a \mid s) da ds \\ &= \mathbb{E}_{s \sim \mu^\pi(s), a \sim \pi_\theta} [q_\pi(S_t, A_t) \nabla \log \pi(A_t \mid S_t)]\end{aligned}$$

Continuous Action Spaces

- ❑ Hard to implement differentiable continuous controllers in many problems
- ❑ **Expectation over both states and actions.** If we can reduce the expectation to only over states, this might simplify gradient estimation

$$\nabla J(\theta) = \int_{\mathcal{S}} \mu^{\pi}(s) \int_{\mathcal{A}} q_{\pi}(s, a) \nabla \pi(a \mid s) da \, ds$$

Deterministic Policies

- ❑ Let the policy $\pi_\theta : \mathcal{S} \rightarrow \mathcal{A}$ be deterministic
- ❑ That is, $\pi_\theta(s)$ will output the action to be taken at state s

- ❑ The performance objective:

$$J(\theta) = \int_{\mathcal{S}} \mu^\pi(s) q_{\pi_\theta}(s, \pi_\theta(s)) ds$$

- ❑ Only one integral now as we have made the policies deterministic

Deterministic Policy Gradient

- ❑ What is the notion of gradient of deterministic policies?
- ❑ Continuous actions spaces allow us to think of change in actions w.r.t the policy parameter

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \int_{\mathcal{S}} \mu^{\pi}(s) \nabla_a q_{\pi}(s, a) \nabla_{\theta} \pi_{\theta}(s) \Big|_{a=\pi_{\theta}(s)} ds \\ &= \mathbb{E}_{s \sim \mu^{\pi}} \left[\nabla_a q_{\pi}(s, a) \nabla_{\theta} \pi_{\theta}(s) \Big|_{a=\pi_{\theta}(s)} ds \right]\end{aligned}$$

- ❑ Expectation only over states
- ❑ Avoided max over actions by moving in the direction of the gradient of q_{π}

DPG Updates

- ❑ For a wide class of stochastic policies, the deterministic policy gradient is a limiting case of the stochastic policy gradient
- ❑ Update equations in actor critic framework:


$$\delta_t = r_t + \gamma \hat{q}(\mathbf{s}_{t+1}, \mathbf{a}_{t+1}, \mathbf{w}) - \hat{q}(\mathbf{s}_t, \mathbf{a}_t, \mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha_{\mathbf{w}} \delta_t \nabla_{\mathbf{w}} \hat{q}(\mathbf{s}_t, \mathbf{a}_t, \mathbf{w})$$

$$\theta \leftarrow \theta + \alpha_{\theta} \nabla_a \hat{q}(\mathbf{s}_t, \mathbf{a}_t, \mathbf{w}) \nabla_{\theta} \pi_{\theta}(\mathbf{s}_t) \big|_{a=\pi_{\theta}(s)}$$

DPG Exploration

- ❑ How do we ensure exploration?
 - ❑ If the environment is stochastic, then not an issue.
 - ❑ Otherwise use off-policy actor critic, where the behaviour policy differs from the estimation policy
- ❑ Requires compatible parametrizations



Deep Deterministic Policy Gradient

DDPG (Deep DPG)

- DDPG combines DPG with DQN

- The algorithm is off-policy. Behaviour policy:

$$\pi'(s) = \pi_{\theta}(s) + N \xrightarrow{\text{Noise}}$$

- Uses replay buffer, improved sample efficiency

- Maintains separate target network parameters θ' , w' and uses soft updates

- Uses Batch Normalization to normalize input state features and minimize covariate shift

DDPG Updates

❑ Critic Update

$$\delta_t = r_t + \gamma \hat{q}(\mathbf{s}_{t+1}, \pi_{\theta'}(\mathbf{s}_{t+1}), \mathbf{w}') - \hat{q}(\mathbf{s}_t, \mathbf{a}_t, \mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha_{\mathbf{w}} \delta_t \nabla_{\mathbf{w}} \hat{q}(\mathbf{s}_t, \mathbf{a}_t, \mathbf{w})$$

Target
networks



❑ Actor Update

$$\theta \leftarrow \theta + \alpha_{\theta} \nabla_a \hat{q}(\mathbf{s}_t, \mathbf{a}_t, \mathbf{w}) \nabla_{\theta} \pi_{\theta}(\mathbf{s}_t) |_{a=\pi_{\theta}(s)}$$

DDPG Algorithm

Randomly initialize critic network $\hat{q}(s, a, \mathbf{w})$ and actor $\pi_\theta(s)$ with weights \mathbf{w} and θ

Initialize target network parameters $\mathbf{w}' \leftarrow \mathbf{w}$, $\theta' \leftarrow \theta$

Initialize replay buffer \mathcal{R}

for $episode=1, M$ do

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation s_1

 for $t=1, T$ do

 Select action $a_t = \pi_\theta(s_t) + \mathcal{N}$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in \mathcal{R}

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from \mathcal{R}

 Set $y_i = r_i + \gamma \hat{q}(s_{i+1}, \pi_{\theta'}(s_{i+1}), \mathbf{w}')$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - \hat{q}(s_i, a_i, \mathbf{w}))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_\theta J \approx \frac{1}{N} \sum_i \nabla_a \hat{q}(s_i, a_i, \mathbf{w}) \nabla_\theta \pi_\theta(s_i) |_{a=\pi_\theta(s)}$$

 Update the target parameters:

$$\mathbf{w}' \leftarrow \tau \mathbf{w} + (1 - \tau) \mathbf{w}'$$

$$\theta' \leftarrow \tau \theta + (1 - \tau) \theta$$

 end

end

} Collect samples

} Update actor & critic

} Update target network