

Soft Actor-Critic

B. Ravindran

Soft Actor-Critic (SAC)

- v1 : "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor", Haarnoja et al
- v2 : "Soft Actor-Critic: Algorithms and Applications", Haarnoja et al

Soft Actor-Critic (SAC)





- v1 : "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor", Haarnoja et al
- v2 : "Soft Actor-Critic: Algorithms and Applications", Haarnoja et al
- Off-policy, model-free RL algorithm
- Uses stochastic policies
- Uses maximum entropy formulation to encourage stability and exploration

Soft Actor-Critic (SAC)

- v1 : "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor", Haarnoja et al
- v2 : "Soft Actor-Critic: Algorithms and Applications", Haarnoja et al
- Off-policy, model-free RL algorithm
- Uses stochastic policies
- Uses maximum entropy formulation to encourage stability and exploration
- Sample-efficient
- Scales to high-dimensional observation/action spaces
- Robust to random seeds, noise etc.

Maximum Entropy RL

- Maximize expected return while acting as randomly as possible
- Agent can capture different modes of optimality to improve robustness against environmental changes







Maximum Entropy?

- Given many distributions that satisfy certain constraints, pick the one that has maximum entropy.

$b \backslash a$	0	1	p_b	$b \backslash a$	0	1	p_b	$b \backslash a$	0	1	p_b
0	p_{11}	p_{12}		0	0.5	0.1		0	0.3	0.2	
1	p_{21}	p_{22}		1	0.1	0.3		1	0.3	0.2	
p_a	0.6	0.4	1	p_a	0.6	0.4	1	p_a	0.6	0.4	1

Maximum Entropy RL

- Maximize expected return while acting as randomly as possible
- Agent can capture different modes of optimality to improve robustness against environmental changes



Maximum Entropy RL

- Entropy of a random variable x

$$H(P) = \mathbb{E}_{x \sim P} [-\log P(x)].$$

- Maximum Entropy RL objective

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t \left(R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot | s_t)) \right) \right]$$

- Here $\alpha > 0$, is the weightage given to the entropy term in the objective. α is commonly referred to as the "temperature"

Maximum Entropy RL

- Define the value function to include the entropy bonuses from every timestep:

$$V^{\pi}(s) = \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t \left(R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot | s_t)) \right) \middle| s_0 = s \right]$$

Maximum Entropy RL

- Define the value function to include the entropy bonuses from every timestep:

$$V^{\pi}(s) = \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t \left(R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot | s_t)) \right) \middle| s_0 = s \right]$$

- Define the action-value function to include the entropy bonuses from every timestep *except the first*:

$$Q^{\pi}(s, a) = \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) + \alpha \sum_{t=1}^{\infty} \gamma^t H(\pi(\cdot | s_t)) \middle| s_0 = s, a_0 = a \right]$$

Maximum Entropy RL

- Thus

$$V^{\pi}(s) = \mathbb{E}_{a \sim \pi} [Q^{\pi}(s, a)] + \alpha H(\pi(\cdot|s))$$

Maximum Entropy RL

- Thus

$$V^\pi(s) = \mathbb{E}_{a \sim \pi} [Q^\pi(s, a)] + \alpha H(\pi(\cdot|s))$$

- Bellman equation

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E}_{\substack{s' \sim P \\ a' \sim \pi}} [R(s, a, s') + \gamma (Q^\pi(s', a') + \alpha H(\pi(\cdot|s')))] \\ &= \mathbb{E}_{s' \sim P} [R(s, a, s') + \gamma V^\pi(s')] . \end{aligned}$$

SAC

- Policy π_{θ}
- Two Q functions Q_{ϕ_1}, Q_{ϕ_2}
- Two target Q functions
- SAC v1 : Temperature α is a hyperparameter
- SAC v2 : Temperature α is learnt

Learning the Q functions

- Both Q-functions are learned with MSBE minimization, by regressing to a single shared target.

$$L(\phi_i, \mathcal{D}) = \mathbb{E}_{(s,a,r,s',d) \sim \mathcal{D}} \left[\left(Q_{\phi_i}(s,a) - y(r,s',d) \right)^2 \right]$$

- The shared target is computed using target Q-networks, which are obtained by polyak averaging the Q-network parameters.
- The shared target makes use of the **clipped double-Q** trick.

$$y(r, s', d) = r + \gamma(1 - d) \left(\min_{j=1,2} Q_{\phi_{\text{targ},j}}(s', \tilde{a}') - \alpha \log \pi_{\theta}(\tilde{a}'|s') \right)$$

- The next-state actions used in the target come from the **current policy** instead of a target policy.

$$\tilde{a}' \sim \pi(\cdot|s')$$

Learning the policy

- Maximize

$$\begin{aligned} V^\pi(s) &= \mathbb{E}_{a \sim \pi} [Q^\pi(s, a)] + \alpha H(\pi(\cdot|s)) \\ &= \mathbb{E}_{a \sim \pi} [Q^\pi(s, a) - \alpha \log \pi(a|s)]. \end{aligned}$$

- Policy is stochastic, therefore actions are sampled.
- To be able to backprop through sampled actions, we use the **reparameterization trick**
 - Policy outputs mean and variance of a Gaussian distribution
 - Add noise
 - Use tanh to squash to $[-1, 1]$

$$\tilde{a}_\theta(s, \xi) = \tanh(\mu_\theta(s) + \sigma_\theta(s) \odot \xi), \quad \xi \sim \mathcal{N}(0, I).$$

- Thus we can rewrite the expectation over actions into an expectation over noise,

$$\mathbb{E}_{a \sim \pi_\theta} [Q^{\pi_\theta}(s, a) - \alpha \log \pi_\theta(a|s)] = \mathbb{E}_{\xi \sim \mathcal{N}} [Q^{\pi_\theta}(s, \tilde{a}_\theta(s, \xi)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s, \xi)|s)]$$

- Thus we can rewrite the expectation over actions into an expectation over noise,

$$\mathbb{E}_{a \sim \pi_\theta} [Q^{\pi_\theta}(s, a) - \alpha \log \pi_\theta(a|s)] = \mathbb{E}_{\xi \sim \mathcal{N}} [Q^{\pi_\theta}(s, \tilde{a}_\theta(s, \xi)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s, \xi)|s)]$$

- Thus the final objective becomes

$$\max_{\theta} \mathbb{E}_{\substack{s \sim \mathcal{D} \\ \xi \sim \mathcal{N}}} \left[\min_{j=1,2} Q_{\phi_j}(s, \tilde{a}_\theta(s, \xi)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s, \xi)|s) \right]$$

Algorithm 1 Soft Actor-Critic

- 1: Input: initial policy parameters θ , Q-function parameters ϕ_1, ϕ_2 , empty replay buffer \mathcal{D}
- 2: Set target parameters equal to main parameters $\phi_{\text{targ},1} \leftarrow \phi_1, \phi_{\text{targ},2} \leftarrow \phi_2$
- 3: **repeat**
- 4: Observe state s and select action $a \sim \pi_\theta(\cdot|s)$
- 5: Execute a in the environment
- 6: Observe next state s' , reward r , and done signal d to indicate whether s' is terminal
- 7: Store (s, a, r, s', d) in replay buffer \mathcal{D}
- 8: If s' is terminal, reset environment state.
- 9: **if** it's time to update **then**
- 10: **for** j in range(however many updates) **do**
- 11: Randomly sample a batch of transitions, $B = \{(s, a, r, s', d)\}$ from \mathcal{D}
- 12: Compute targets for the Q functions:

$$y(r, s', d) = r + \gamma(1 - d) \left(\min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

- 13: Update Q-functions by one step of gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$

- 14: Update policy by one step of gradient ascent using

$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} \left(\min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_\theta(s)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s)|s) \right),$$

where $\tilde{a}_\theta(s)$ is a sample from $\pi_\theta(\cdot|s)$ which is differentiable wrt θ via the reparametrization trick.

- 15: Update target networks with


$$\phi_{\text{targ},i} \leftarrow \rho \phi_{\text{targ},i} + (1 - \rho) \phi_i \quad \text{for } i = 1, 2$$

- 16: **end for**
 - 17: **end if**
 - 18: **until** convergence
-


SAC v1

- "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor", Haarnoja et al
- Temperature α is a hyperparameter
- Tasks
 - A range of continuous control tasks from the OpenAI gym benchmark suite
 - The easier tasks can be solved by a wide range of different algorithms. The more complex benchmarks, such as the 21-dimensional Humanoid (rllab) are exceptionally difficult to solve with off-policy algorithms.
- Baselines:
 - DDPG, SQL, PPO, TD3 (concurrent)
 - TD3 is an extension to DDPG that first applied the double Q-learning trick to continuous control along with other improvements.


SAC v1 : Experimental Results




(a) Hopper-v1




(b) Walker2d-v1




(c) HalfCheetah-v1



(d) Ant-v1




(e) Humanoid-v1



(f) Humanoid (rllab)


SAC v1 : Ablation Study

- How does the stochasticity of the policy and entropy maximization affect the performance?
- Comparison with a deterministic variant of SAC that does not maximize the entropy and that closely resembles DDPG




SAC v1 :


Hyperparameter Sensitivity



(a) Evaluation



(b) Reward Scale



(c) Target Smoothing Coefficient (τ)

Limitation of SAC v1

- SAC v1 is brittle to the choice of hyperparameter α (temperature) that controls exploration
 - Solution --> Automatic temperature tuning!

SAC v2

- "Soft Actor-Critic: Algorithms and Applications", Haarnoja et al
- Temperature α is learnt
- Shows results on simulated tasks from OpenAI gym, RL Lab as well as real-world tasks such as locomotion for a quadrupedal robot and robotic manipulation with a dexterous hand

Experimental Results: RL Lab

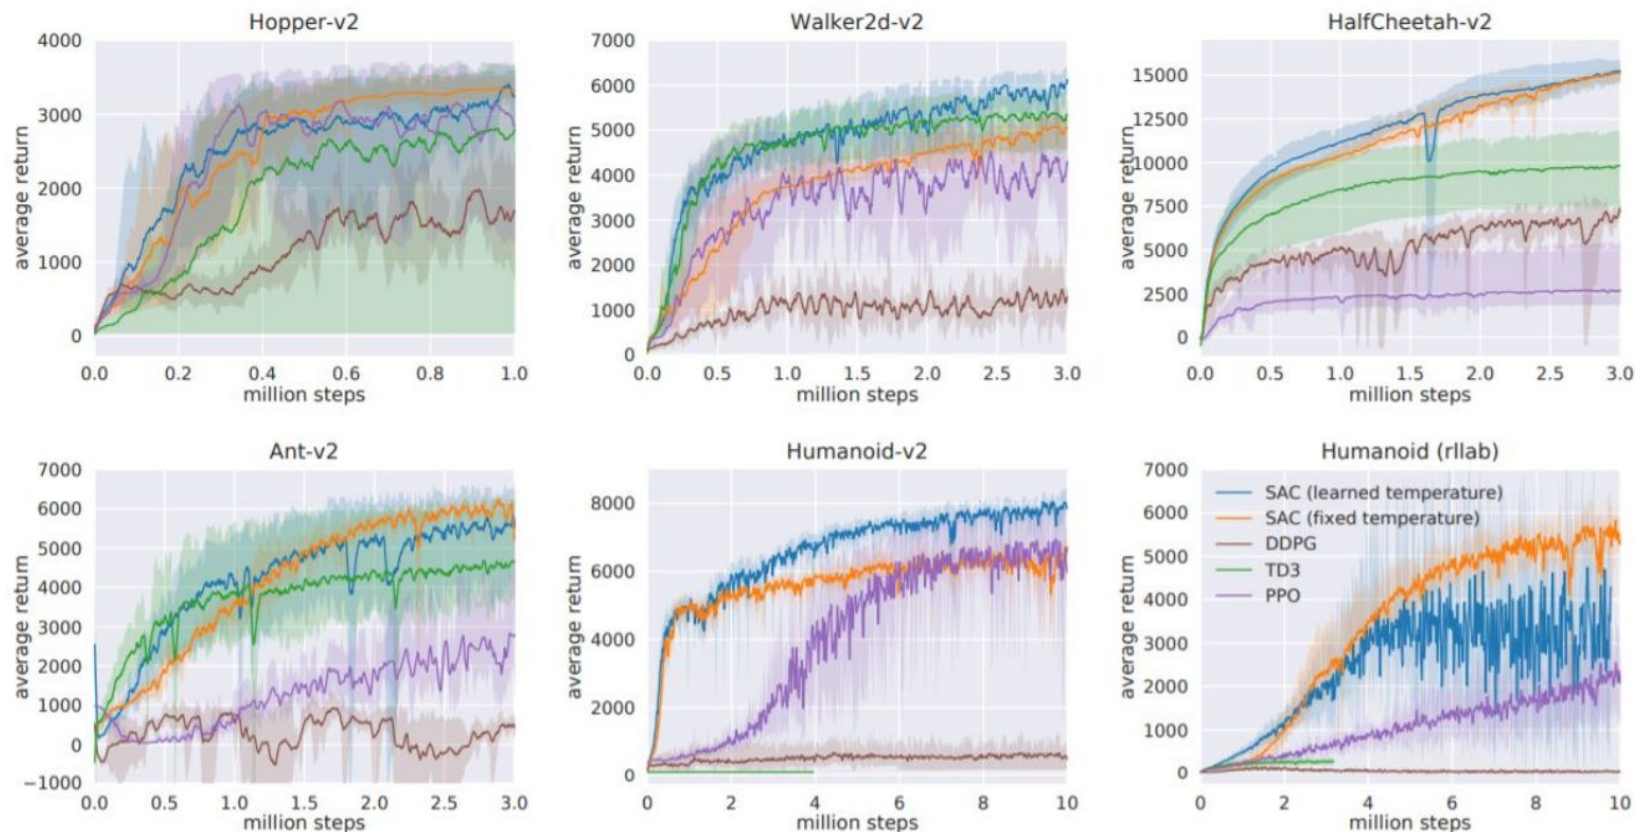



Figure 1: Training curves on continuous control benchmarks. Soft actor-critic (blue and yellow) performs consistently across all tasks and outperforming both on-policy and off-policy methods in the most challenging tasks.

Real World Robots

Quadrupedal Robot Locomotion




Real World Robots

- Dexterous hand manipulation
 - 20 hour end-to-end learning
 - Valve position as input: SAC 3 hours vs. PPO 7.4 hours



Limitations/Open Issues

- Lack of experiments on hard-exploration problems
- High-variance due to automatic temperature tuning



Recap: SAC

- An off-policy, model-free maximum entropy deep RL algorithm
 - Sample-efficient
 - Scales to high-dimensional observation/action spaces
 - Robust to random seeds, noise etc.
- SAC outperforms SOTA model-free deep RL methods, including DDPG, PPO in terms of the policy's optimality, sample complexity and robustness.

Thank you !!